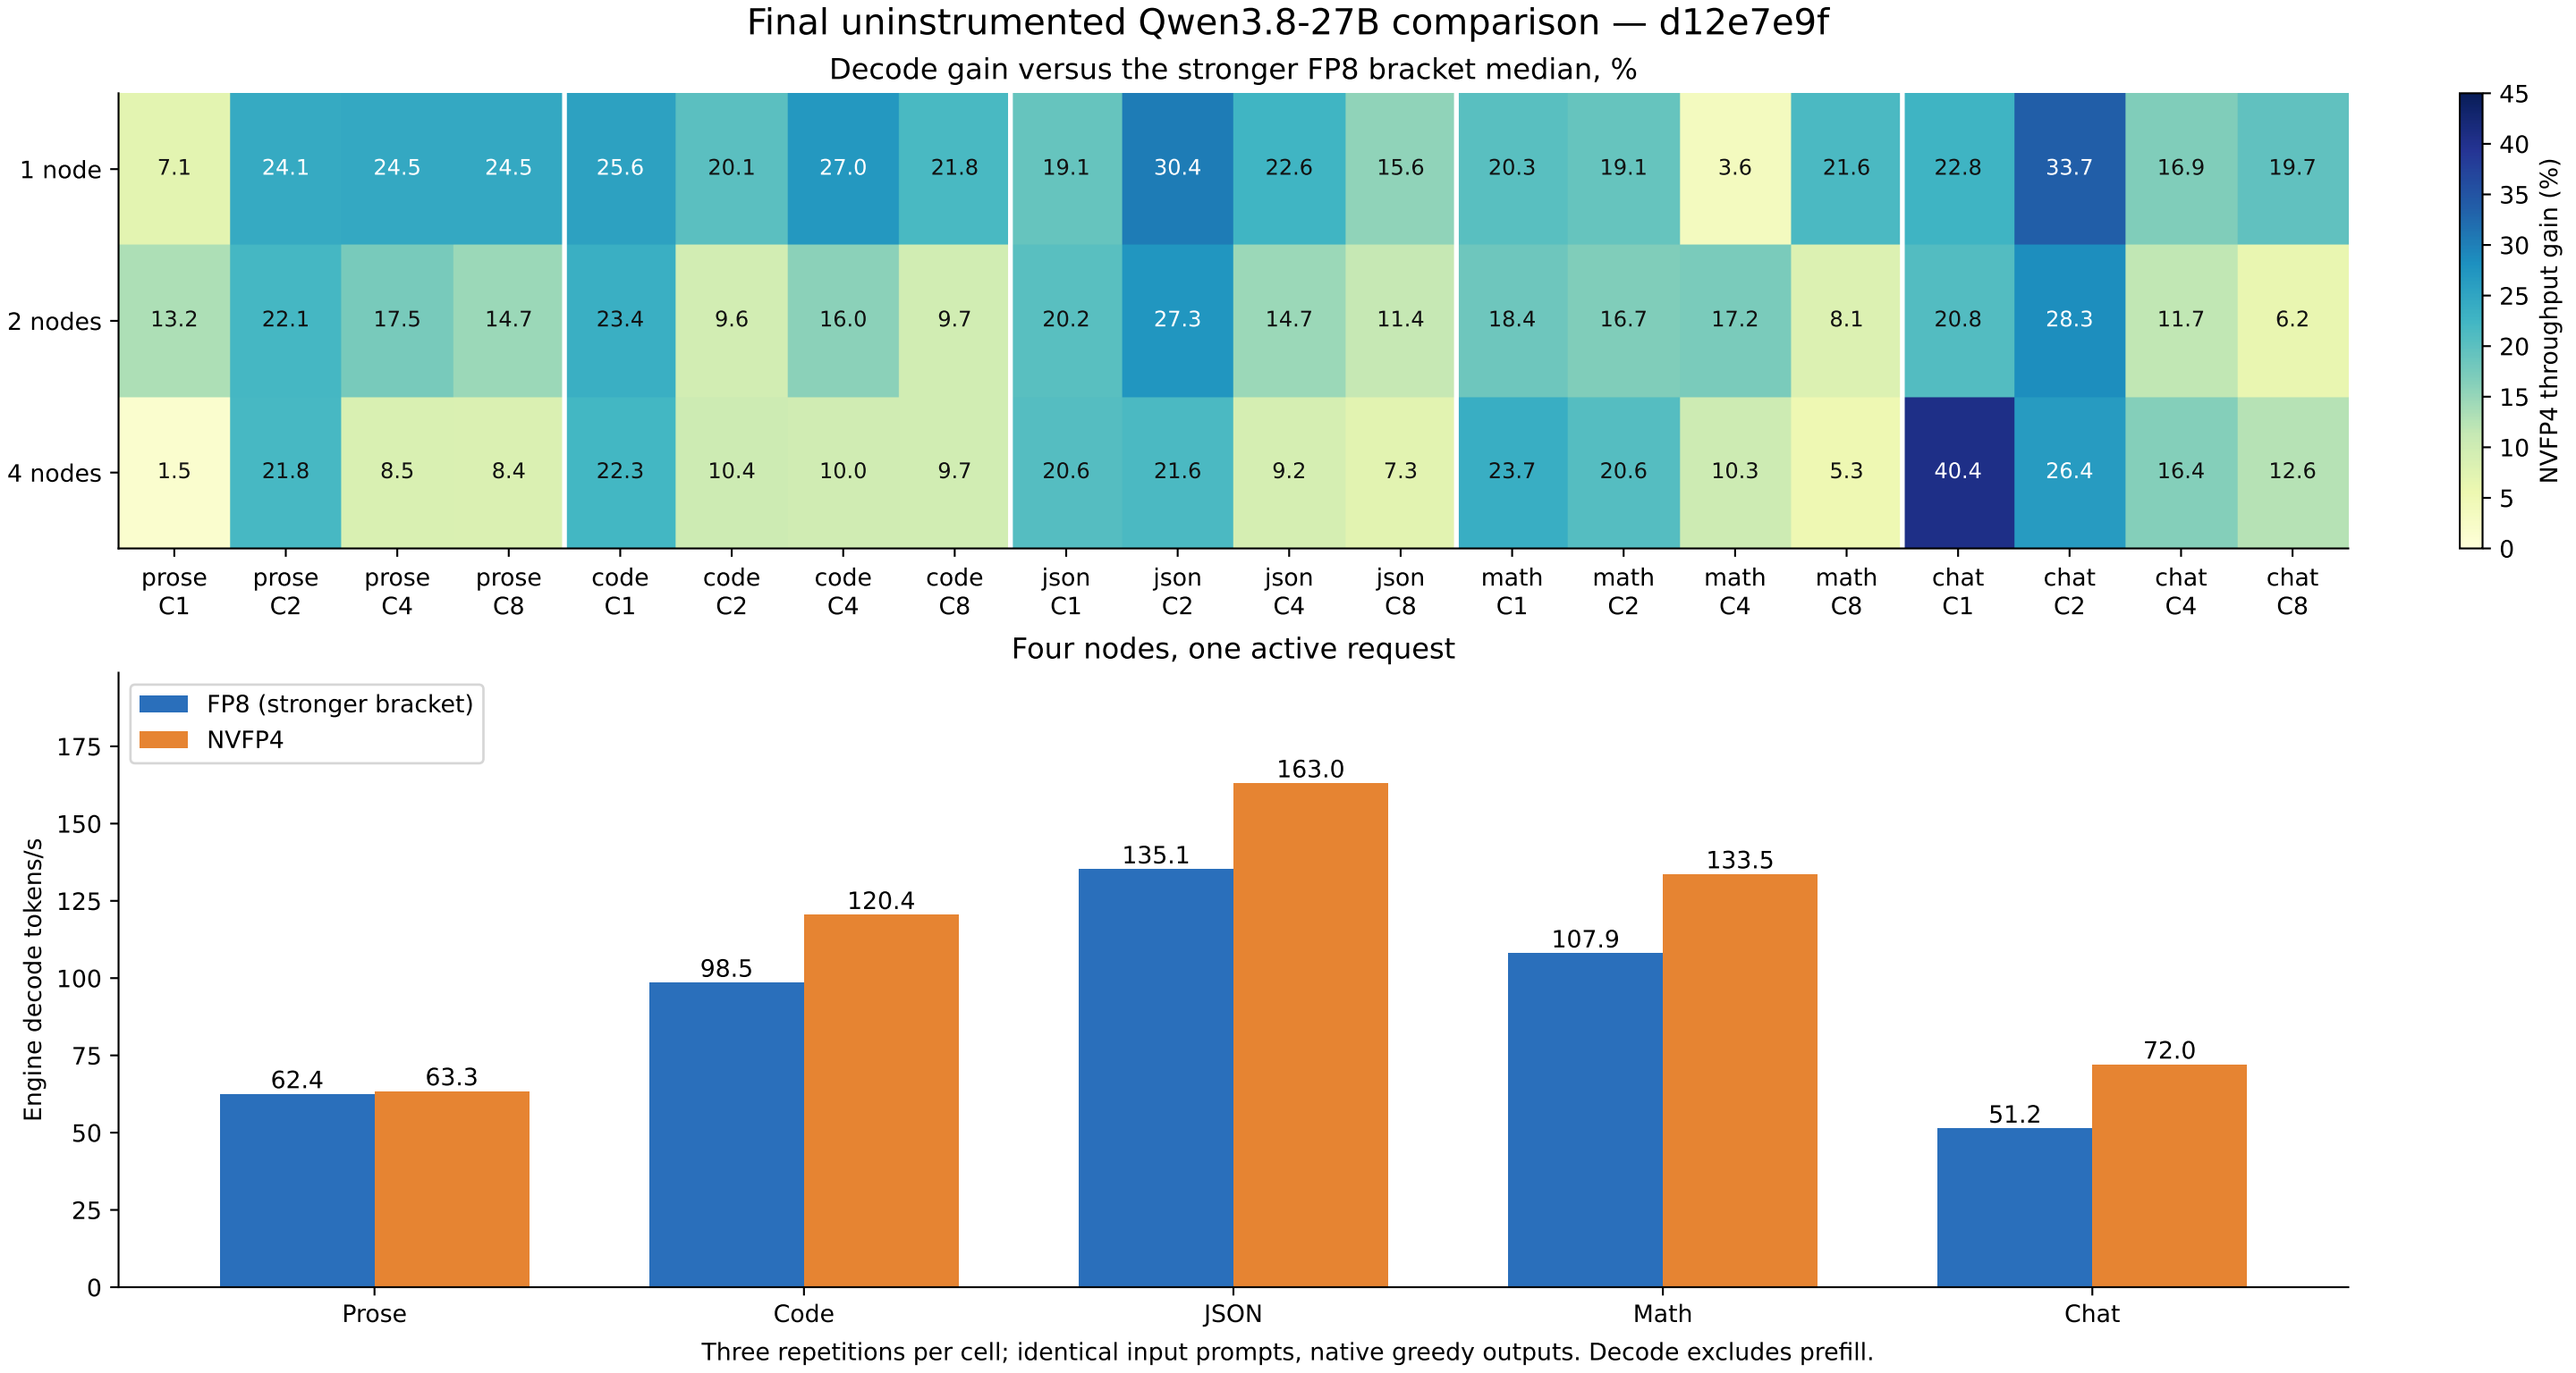
Final uninstrumented Qwen3.8-27B comparison — d12e7e9f
Decode gain versus the stronger FP8 bracket median, %
1 node
2 nodes
4 nodes
7.1 24.1 24.5 24.5 25.6 20.1 27.0 21.8 19.1 30.4 22.6 15.6 20.3 19.1 3.6 21.6 22.8 33.7 16.9 19.7
13.2 22.1 17.5 14.7 23.4 9.6 16.0 9.7 20.2 27.3 14.7 11.4 18.4 16.7 17.2 8.1 20.8 28.3 11.7 6.2
1.5 21.8 8.5 8.4 22.3 10.4 10.0 9.7 20.6 21.6 9.2 7.3 23.7 20.6 10.3 5.3 40.4 26.4 16.4 12.6
prose C1
prose C2
prose C4
prose C8
code C1
code C2
code C4
code C8
json C1
json C2
json C4
json C8
math C1
math C2
math C4
math C8
chat C1
chat C2
chat C4
chat C8
NVFP4 throughput gain (%)
0
5
10
15
20
25
30
35
40
45
Four nodes, one active request
FP8 (stronger bracket)
NVFP4
Engine decode tokens/s
0
25
50
75
100
125
150
175
62.4
63.3
98.5
120.4
135.1
163.0
107.9
133.5
51.2
72.0
Prose
Code
JSON
Math
Chat
Three repetitions per cell; identical input prompts, native greedy outputs. Decode excludes prefill.