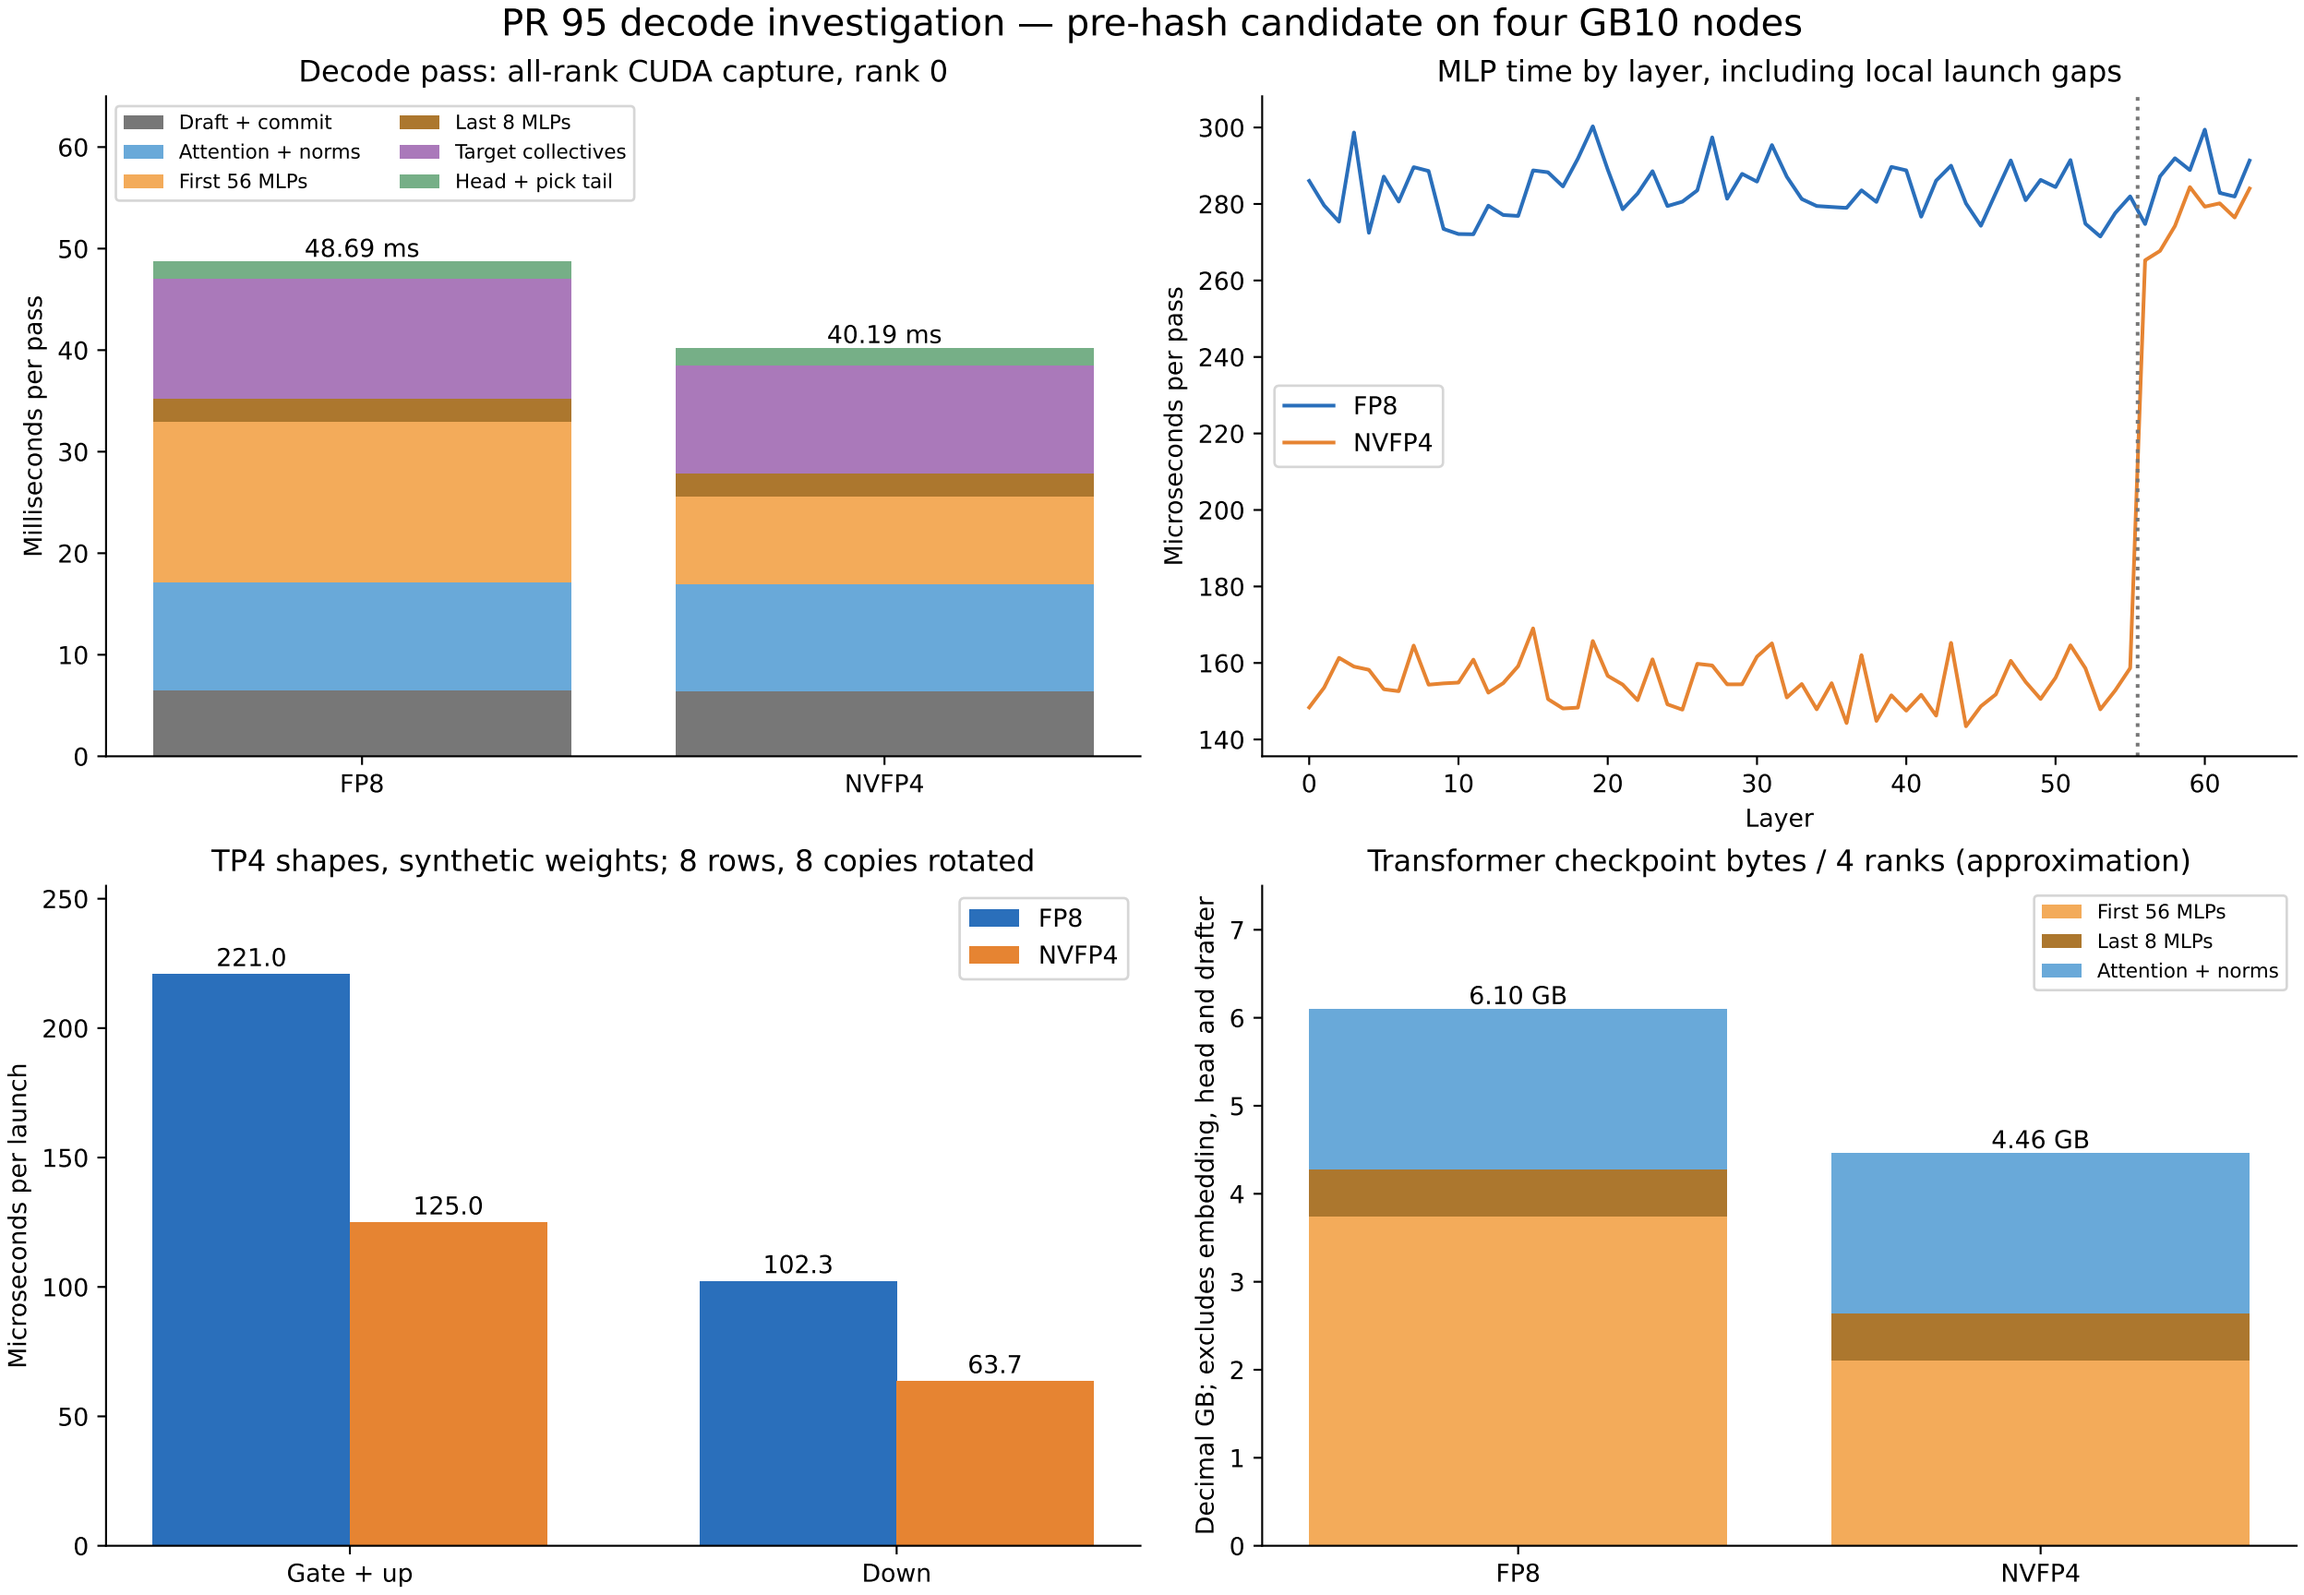
PR 95 decode investigation — pre-hash candidate on four GB10 nodes
Decode pass: all-rank CUDA capture, rank 0
Draft + commit
Attention + norms
First 56 MLPs
Last 8 MLPs
Target collectives
Head + pick tail
48.69 ms
40.19 ms
Milliseconds per pass
FP8
NVFP4
MLP time by layer, including local launch gaps
Microseconds per pass
Layer
TP4 shapes, synthetic weights; 8 rows, 8 copies rotated
221.0
125.0
102.3
63.7
Microseconds per launch
Gate + up
Down
Transformer checkpoint bytes / 4 ranks (approximation)
6.10 GB
4.46 GB
Decimal GB; excludes embedding, head and drafter